TR1: Stability Study of Ceph RBD

Xing Lin
xinglin@cs.utah.edu
University of Utah

03/08/2013

1 Introduction

This document presents the results about the performance stability of Ceph Rados Block Device (RBD).

2 Experiment Setup

I did my experiments in the Emulab network testbed, hosted by the FLUX group at the University of Utah. I used 4 d820 nodes as the Ceph cluster and another d820 node as the client machine. Each d820 node has 6*600 GB SCSI disks so in total we have 4*6 (24) disks. We used xfs as the file system for each osd. The journal is set to be 10 GB. To simulate the situation where the journal disk is hosted by a SSD, I put the journal disk into a tmpfs file (by specifying osd journal = /dev/shm/journal/$name-journal in /etc/ceph/ceph.conf).

The fio tool is used to generate synthetic workloads.

3 Results

3.1 Sequential Write

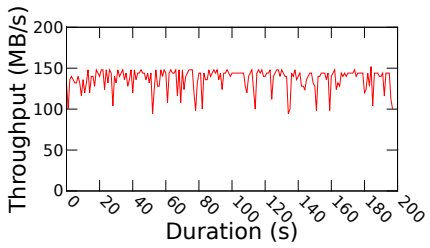

(a) IODepth=1

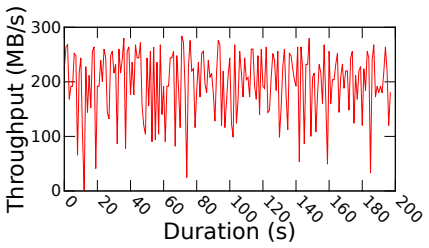

(b) IODepth=2

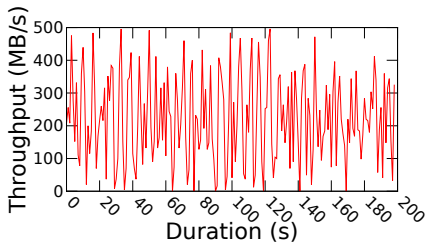

(c) IODepth=4

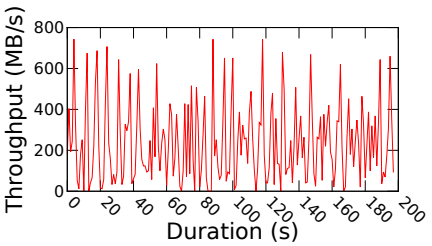

(d) IODepth=8

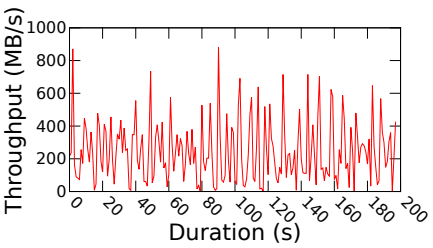

(e) IODepth=16

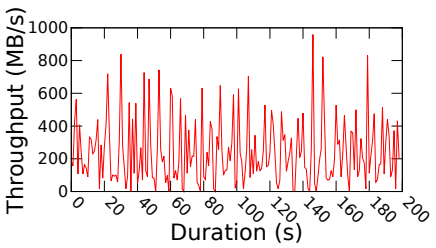

(f) IODepth=32

Figure 1: Instant Throughputs of a 4M Sequential Write Workload