

# TR4: Scalability Study of XFS

Xing Lin  
xinglin@cs.utah.edu  
University of Utah

03/08/2013

## 1 Introduction

This document presents the results about the performance scalability of XFS.

## 2 Experiment Setup

I did my experiments with a single d820 machines in the Emulab network testbed, hosted by the FLUX group at the University of Utah. An xfs was created on /dev/sdb and fio was used to generated synthetic workloads.

For scalability, I mainly looked into three aspects: IO depth (parallel writes to a single file), block size, and parallel writes to many files.

| tool       | version                           |
|------------|-----------------------------------|
| processors | 4 × 2.2 GHz 64-bit 8-Core E5-4620 |
| Memory     | 128 GB DDR3 RAM                   |
| Disk       | 10K RPM SCSI MBF2600RC (TOSHIBA)  |
| OS         | Ubuntu12.04                       |
| fio        | 2.0.14                            |

Table 1: Tools

## 3 Results

### 3.1 IO Depth

I first looked into how iodepth affects the throughput of a single sequential workload. The parameters are presented in Table ?? and the results are presented in Figure ?? and ?? .

| block size | duration | directio | ioengine |
|------------|----------|----------|----------|
| 4 MB/4 KB  | 60 s     | 1        | libaio   |

Table 2: Sequential workload parameters




Figure 1: The average bandwidth of a SW workload, varying IO depth




Figure 2: The average bandwidth of a SR workload, varying IO depth

### 3.2 Block Size

In this measurement, I fixed iodepth=1 and varied the block size. The parameters of workloads used in this set of experiments are presented in the Table ?? and results are presented in Figure ??.

### 3.3 Parallel Writers/Readers

In this set of experiments, iodepth was fixed at 1 and the number of files that were concurrently accessed was varied. The parameters about the workloads are presents in Table ?? and the results are presented in Figure ?? and ??.

| duration | directio | ioengine | iodepth |
|----------|----------|----------|---------|
| 60 s     | 1        | libaio   | 1       |

Table 3: Sequential workload parameters




Figure 3: Average throughputs of sequential workload, varying the block size




Figure 4: Average throughputs of a SW workload




Figure 5: Average throughputs of a SR workload

## 4 Discussions

Figure ?? and ?? shows bandwidths general increase as iodepth is increased for block size=4 KB while it does not for block size=4 MB. That can be explained in that as the iodepth

| block size | duration | directio | ioengine | iodepth |
|------------|----------|----------|----------|---------|
| 4 MB/4 KB  | 60 s     | 1        | libaio   | 1       |

Table 4: Sequential workload parameters

increases, there are greater opportunities to merge many small sequential requests into larger ones. Figure ??(a) and ??(a) shows an unexpected bandwidth drop at iodepth=32, for block size = 4 KB.

Figure ?? shows that block size starting from 32 KB can reach the maximum throughput xfs can provide for a single sequential access stream.

Figure ?? and ?? shows that as we increases the number of concurrent files to read from/write into, the aggregated throughput drops dramatically. This effect is more significant for small block sizes than large block sizes. In particular, we should avoid co-locate **sequential write workloads with small block sizes** at the same disk.

## 5 TODO

1. extended with random workloads